

# Thread Scheduler Efficiency Improvements for Multiprocessor Multicore Systems

Daniel C. Frazier  
Division of Science and Mathematics  
University of Minnesota, Morris  
Morris, Minnesota, USA 56267  
frazi177@morris.umn.edu

## ABSTRACT

[Abstract contents]

## Keywords

Scheduling; thread migration; multicore; multiprocessing;  
lock contention; last-level cache misses

## 1. INTRODUCTION

[Introduction contents]

## 2. BACKGROUND

In this Section we will broadly establish how threading, scheduling and caching works. We will establish causes of cache misses in highly parallel programs mediated by the current implementation of thread scheduler written for Linux.

### 2.1 Threads and Scheduling

Using any modern computer, there is an expectation that the operating system is running at all times (largely in the background) and that multiple different user programs and system programs should be able to run concurrently. Modern computer programs also often need to run more than one independent task at one time. This can be accomplished by employing *threads*, or equivalently, making the program *multithreaded*. Programs that involve long independent computations or programs with a graphical interface often benefit from employing threads. Threads are tied to the processes that spawn them. A process always has at least one thread. Processes are typically independent of each other while threads exist within a process. Threads within a process share address space while different processes have distinct address spaces. Context switching is typically faster for threads than processes because most of the state information between threads in a process is the same [5]. Context switching within a CPU is the process of storing and restoring the state of a process or thread so that execution can be paused or resumed. [4].

The part of the operating system responsible for managing the CPU runtime that each of these processes and their respective threads is called the scheduler. New threads and

processes are added to the scheduler when they are made [2].

### 2.2 Completely Fair Scheduler (CFS)

Scheduler implementations vary per operating system. The scheduler used on Linux is called the Completely Fair Scheduler (CFS.) The CFS creates timeslices. CPU cycles are distributed between the amount of available threads.

[2]

Description of CFS, linux implementation of Weighted Fair Queueing (WFQ) algorithm and how it originally intended for single-core use [2].

### 2.3 Cache on NUMA Systems

*Non-uniform memory access (NUMA)*: There are many levels of cache and they exist in a hierarchy. L1 Cache exists on a processor or on a small group of processors. L2 (and higher) Cache exists on increasing sizes of groupings of processors. Cache misses happen and they reduce performance substantially.

### 2.4 Cache Locality

It's faster to access data in L1 Cache than L2 and above. Allocating related threads to processors such that they share cache improves locality of cache lookups and are less likely to have cache misses.

### 2.5 Faults of the CFS

CFS modifications were made to work for multiprocessor systems and it is buggy and leaves threads waiting while a processor is idle for short, previously undetected amounts of time. Give stats for time idle [2]. Kumar et al. defines LLC miss rate as the last level cache misses per thousand instructions (MPKI) [1].

## 3. METHODS

So far we've seen how threads are being mismanaged by the CFS and how they impact performance. Next, we'll look at some recent developments that reduce lock contention.

### 3.1 Shuffler and Jumbler

The CFS doesn't differentiate between threads of a single-threaded program versus the threads of a multithreaded program. This prevents the scheduler from using that metadata in its thread distribution mechanism. The following thread schedulers improve upon this.

#### 3.1.1 Shuffler

Performance of multithreaded applications on multicore multiprocessor systems with high lock contention is dependant on the distribution of those threads across processors. The Shuffling approach is for the scheduling algorithm to take into account what threads are contending for locks on what processors and migrate threads to share processors. The Shuffling Framework does this by the following. First we find the expected arrival time of locks on threads. Then we sort these threads by their expected arrival times and group them in as many groups as there are processors. Then we distribute these groups of threads to their own respective processors. See Algorithm 1. The migration of threads between processors is costly, but the thread that is waiting during that time is not doing any useful work anyway. In addition, contending threads that are co-located in one processor can share data much faster and avoid LLC misses. For these reasons it is preferable to migrate the whole thread rather than it's data and the lock. This will be shown in the Performance section [1].

For the first step of the algorithm you must find the expected arrival time of threads. The *lock time* of a thread is measured by the percent of time it spends waiting for locks. This is only done for threads in user space. There exists a daemon thread that contains a data structure that maps threads to their lock times and processor ids. For the monitor we must choose a rate to sample lock times and a rate to perform thread migration (shuffling). We use `prstat(1)` to monitor lock times. Kumar et al found that finer sampling rates allow for detailed monitoring but also more overhead. They found that for sampling rates less than 200 ms, the overhead was significantly higher. For a lock sampling rate of 200 ms the process of sampling took less than 1% system time. The Shuffling interval was chosen by experimentation. We tested various shuffling intervals on 20 programs and chose 500 ms.

On an iteration of the grouping-forming procedure (every 200 ms), the daemon checks the total amount of time that was spent resolving locks on each thread and if that time exceeds a preset limit, then groups are formed again.

On an iteration of the shuffling procedure (every 500 ms), shuffling checks to make sure that if any threads aren't on processors that they were grouped to, they are migrated. Threads that are already on the processor that they are assigned do not migrate. If nothing substantially different changed in how threads interacted, the shuffling step is effectively skipped [?].

```

input : N: Number of threads; C: Number of
        Sockets.
repeat
  i. Monitor Threads – sample lock times of N
    threads.
  if lock times exceed threshold then
    ii. Form Thread Groups – sort threads
        according to lock times and divide them into
        C groups.
    iii. Perform Shuffling – shuffle threads to
        establish newly computed thread groups.
  end
until application terminates;
Algorithm 1: The Shuffling Framework.

```

### 3.1.2 Jumbler

### 3.1.3 Shuffler and Jumbler Performance

In Kumar et al we studied the lock times of 33 programs on a 64-core, 4-processor machine running Oracle Solaris 11. We identified 20 of those programs that experienced overall high lock times and used those programs to compare shuffling versus the standard scheduler.

## 3.2 FLSCHED for Xeon Phi Manycore Processor

[Body text]

### 3.2.1 Lockless Thread Scheduler

[2, 3] [1]

### 3.2.2 FLSCHED Performance

[Body text]

## 4. CONCLUSIONS

[Conclusion text]

## Acknowledgments

## 5. REFERENCES

- [1] K. Kumar, P. Rajiv, G. Laxmi, and N. Bhuyan. Shuffling: A framework for lock contention aware thread scheduling for multicore multiprocessor systems. In *2014 23rd International Conference on Parallel Architecture and Compilation Techniques (PACT)*, pages 289–300, Aug 2014.
- [2] J.-P. Lozi, B. Lepers, J. Funston, F. Gaud, V. Quéma, and A. Fedorova. The linux scheduler: A decade of wasted cores. In *Proceedings of the Eleventh European Conference on Computer Systems, EuroSys '16*, pages 1:1–1:16, New York, NY, USA, 2016. ACM.
- [3] U. B. Nisar, M. Aleem, M. A. Iqbal, and N. S. Vo. Jumbler: A lock-contention aware thread scheduler for multi-core parallel machines. In *2017 International Conference on Recent Advances in Signal Processing, Telecommunications Computing (SigTelCom)*, pages 77–81, Jan 2017.
- [4] Wikipedia. Context switch — Wikipedia, The Free Encyclopedia, 2017. [Online; accessed 11-October-2017].
- [5] Wikipedia. Thread (computing) — Wikipedia, The Free Encyclopedia, 2017. [Online; accessed 10-October-2017].