

# Thread Scheduler Efficiency Improvements for Multicore Systems

Daniel Collin Frazier

Division of Science and Mathematics  
University of Minnesota, Morris  
Morris, Minnesota, USA

18 November 2017  
UMM, Minnesota

# Introduction

- *Thread scheduler*: system component that manages the processing programs receive
- Always running, so it must be efficient
- Pre-2000 single-core era, scheduling was easy
- Led majority of Linux community to believe problem solved

*“...not very many things ... have aged as well as the scheduler. Which is just another proof that scheduling is easy.”*

Linus, Torvals, 2001 [3]

# Introduction

- Popular hardware changed rapidly throughout the 2000s
- Increasing affordability and adoption of multicore systems
- Hardware changes complicated thread scheduler implementation
- Complexity led to bugs that have been present for a decade

# A Decade of Wasted Cores

- Lozi et al. found four bugs in Linux thread scheduler, fixed them [3]
- Previously undetected, required the development of new tools to notice them



<https://goo.gl/3wsfVU>

## A Decade of Wasted Cores

- Lozi et al. compared performance benchmarks ran on **buggy** and **fixed** Linux scheduler implementations
- Below are average performance improvements

Bug title	Improvement
The Scheduling Group Construction bug	5.96x
The Group Imbalance bug	1.05x
The Overload-on-Wakeup bug	1.13x
The Missing Scheduling Domains bug	29.68x

# Outline

Concepts

Thread Scheduling on Linux

Two New Schedulers

Conclusion

# Outline

## Concepts

- Threads

- Synchronicity and Locks

- Thread State and Cache

Thread Scheduling on Linux

Two New Schedulers

Conclusion




# Processors

- Responsible for executing code
- Contain a number of cores:
  - *Single-core processor* (one processing unit)
  - *Multicore processor* (two or more processing units)
  - *Manycore processor* (20 or more processing units)
- Multiple cores allows processor to perform multiple tasks concurrently on each core


# Multithreading Example

- Imagine you're using photoshop, but assume one thread
- Say you load a large image and perform an expensive filter operation



# Threading

- *Threads* allow programs to run multiple independent tasks concurrently
- Useful for programs:
  - with long, mostly-independent computations
  - with a graphical interface



Example GUI Program.

**Three** threads are created within **one** process

What if I ask you all a question right now?

# Synchronicity and Locks

- Control achieved by employing locks
- *Locks* secure objects or data shared between threads so that only one thread can read and write to it at one time
- When a thread *locks* a lock it **acquires** the lock
- When a thread *unlocks* a lock it **releases** the lock

# Process and Thread State

- Process State

Resources shared amongst its multiple threads

- Thread State

Scheduler uses this information to pause and resume a thread's execution


- Note: Process states are much heavier than thread states

# Context Switching

- The scheduler *switches* active threads on cores by saving and restoring thread and processor state information.
- These switches are called *context switches*

# Cache

- Local copy of data designed for fast retrieval
- Hierarchical structure
- Placement relative to core:
  - on
  - inside of
  - outside




**Figure:** Distance of various forms of memory from CPU



# Cache

- *Locality*: Speed of memory read and writes decrease as distance from CPU increases
- Cache is the fastest form of memory
- *Cache coherence*: Any changes to memory shared by two caches must propagate to the other to maintain correctness



**Figure:** Distance of various forms of memory from CPU

# Outline

Concepts

Thread Scheduling on Linux  
Completely Fair Scheduler

Two New Schedulers


Conclusion

# Completely Fair Scheduler (CFS)

- Default Linux thread scheduler (there are others)
- Handles which threads are executed at what times on this core
- Spend a *fair* amount of runtime on all threads
- Designed with responsiveness and fairness in mind.

# Single-core Completely Fair Scheduler (CFS)

- Runs on one core
- Ensure all threads run *at least once* within arbitrary interval of CPU cycles
- Distribute *timeslices* (max CPU cycles) among threads
- Threads with higher priority (weights) get larger timeslices



# CFS Runqueue

- Data structure containing threads
- Priority queue: sorts threads by number of cycles consumed in current interval
- When thread reaches its maximum cycles, preempted

## Runqueues on Multiple Cores

- Process states heavier than thread states, so context switches between threads of different processes are more expensive
- If cores shared a runqueue, access and changes need to be synchronous and cache-coherent
- Would slow the system to crawl
- So each core has its own runqueue and threads
- Load on each of the core's runqueues must stay balanced
- CFS periodically runs a load-balancing algorithm

# Outline

Concepts

Thread Scheduling on Linux

**Two New Schedulers**

**Shuffler**

**FLSCHED**

Conclusion

# Shuffler and FLSCHEd

- Both schedulers aim to solve the same problem, but for different architectures
- **Problem:** Adding more threads to certain parallel computing applications on CFS makes the application operate slower rather than faster!
- Architectures:

Shuffler      →    *multiprocessor multicore*

FLSCHEd     →    *single-chip manycore processor*



# Shuffler

- Researchers Kumar et al. measured lock times of massively parallel applications
- *Lock times*: amount of time process spends waiting for locks
- Found that massively parallel shared-memory programs experienced high lock times

# Lock Contention

- When two threads repeatedly contend for one lock, both threads are frequently waiting for each other to release
- If the two threads are located on separate processors, this problem is compounded by reduced locality
- Further, when both of the threads repeatedly modify the data corresponding to their lock, the cache of both processors must continue to update each other
- High *lock contention*

# Shuffler

- CFS not mindful of lock contention or parent processes when choosing cores for threads
- Kumar et al. wanted to create a scheduler that did!
- Used Solaris scheduler as base
- **Strategy:** Migrate threads whose locks are contending near each other
- How do you determine which threads' locks are contending?
- Contending threads have similar lock acquisition times

**input** : N: Number of threads;  
C: Number of Processors.

**repeat**

**i. Monitor Threads** – sample lock times of N threads.

**if** *lock times exceed threshold* **then**

**ii. Form Thread Groups** – sort threads according to  
        lock times and divide them into C groups.

**iii. Perform Shuffling** – shuffle threads to establish  
        newly computed thread groups.

**end**

**until** *application terminates*;

## Shuffler Performance

- Kumar et al. compared the efficiency of Shuffler vs Solaris scheduler
- Used programs from four benchmarks to gather data

Program	% Improvement
<i>BT</i>	54.1%
<i>SC</i>	29.0%
<i>RX</i>	19.0%
<i>JB</i>	14.0%
<i>OC</i>	13.4%
<i>AL</i>	13.2%
<i>AS</i>	13.0%
<i>PB</i>	13.0%
<i>VL</i>	12.8%
<i>FS</i>	12.0%

Program	% Improvement
<i>FM</i>	10.7%
<i>AM</i>	9.3%
<i>GL</i>	9.1%
<i>EQ</i>	9.0%
<i>MG</i>	8.8%
<i>FA</i>	6.0%
<i>WW</i>	5.2%
<i>SM</i>	4.7%
<i>GA</i>	4.0%
<i>RT</i>	4.0%

# FLSCHED: The Lockless Monster

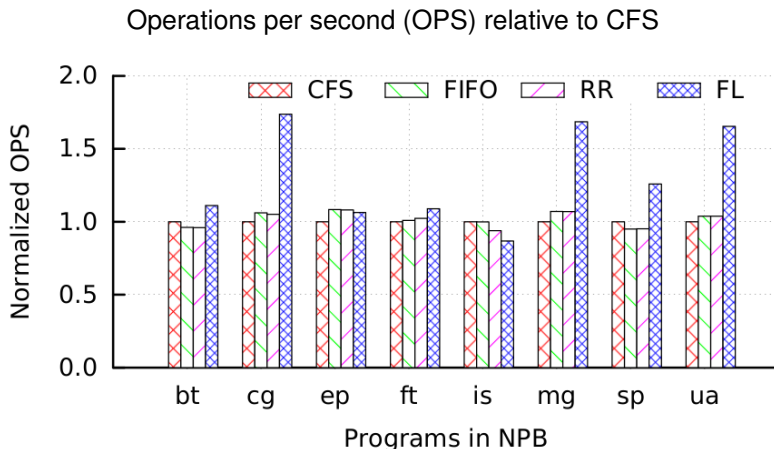
- Designed by Jo et al. with manycore processors in mind, particularly the Xeon Phi
- The Xeon and Xeon Phi have 24 to 76 cores.
- One processor, so cache looks different than system that would use Shuffler
- With such parallelism, small pauses significantly reduce efficiency
- In the CFS, pauses come from locks necessitated by its features and requirements

# One requirement to rule them all: EFFICIENCY!

- FLSCHED Improves efficiency by removing all locks from the scheduler implementation
- Gutted requirements and features of CFS and simplified
- Requirements they removed were **Fairness** and **Responsiveness**
- Context switches requests delayed to reduce chance another thread steals the core in hope thread reactivates
- Threads never forcefully preempt, instead join runqueue with high priority
- Removed scheduler statistics reporting capabilities

# FLSCHED Performance

- Like in Shuffler, used a benchmark of various problems





# Outline

Concepts

Thread Scheduling on Linux

Two New Schedulers

Conclusion

# Conclusion

- Thread scheduling is an important problem and becomes more relevant as number of cores increase
- System architecture can have surprising complexity in its effect on efficiency
- CFS tries to be the go-to scheduler for all problems, but can't
- Does well, but when you need some extra push there are powerful alternatives available

# Thanks!

Thank you for your time and attention!

# Questions?

# References



Jo, Heeseung and Kang, Woonhak and Min, Changwoo and Kim, Taesoo.

FLSched: A lockless and lightweight approach to OS scheduler for Xeon Phi.

In Proceedings of the 8th Asia-Pacific Workshop on Systems 3 *APSys '17*, pages 8:1–8:8, Mumbai, India, 2017. ACM.



K. Kumar and P. Rajiv and G. Laxmi and N. Bhuyan

Shuffling: A framework for lock contention aware thread scheduling for multicore multiprocessor systems

In 2014 23rd International Conference on Parallel Architecture and Compilation Techniques 3 *PACT* , pages 289–300, 2014.



Lozi, Jean-Pierre and Lepers, Baptiste and Funston, Justin and Gaud, Fabien and Quéma, Vivien and Fedorova, Alexandre